

A DETAILED DERIVATION OF I/O LOWER BOUND PROOFS

Detailed Derivation Of I/O Lower Bound Proofs

A.1 Computational Intensity

We start by proving Lemma 2:

Lemma 2. *If $|H_{max}|$ can be expressed as a closed-form function of X , that is there exists some function χ such that $|H_{max}| = \chi(X)$, then the lower bound on Q may be expressed as*

$$Q \geq n \frac{(X_0 - M)}{\chi(X_0)},$$

where $X_0 = \arg \min_X \rho = \arg \min_X \frac{\chi(X)}{X-M}$.

Intuition. $\chi(X)$ expresses computation “volume”, while X is its input “surface”. The term $X - M$ bounds the required communication and it comes from the fact that not all inputs have to be loaded (at most M of them can be reused). X_0 corresponds to the situation where the ratio of this “volume” to the required communication is minimized (corresponding to a highest lower bound).

PROOF. Note that Lemma 1 is valid for any X_c (i.e., for any X_c , it gives a valid lower bound). Yet, these bounds are not necessarily tight. As we want to find tight I/O lower bounds, we need to maximize the lower bound. X_0 by definition minimizes ρ ; thus, it maximizes the bound. Lemma 2 then follows directly from Lemma 1 by substituting $\rho = \frac{\chi(X_0)}{X_0 - M}$. \square

Note. If function $\chi(X)$ is differentiable and has a global minimum, we can find X_0 by, e.g., solving the equation $\frac{d \chi(X)}{dX} = 0$. The key limitation is that it is not always possible to find χ , that is, to express $|H_{max}|$ solely as a function of X . However, for many linear algebra kernels $\chi(X)$ exists. Furthermore, one can relax this problem preserving the correctness of the lower bound, that is, by finding a function $\hat{\chi} : \forall X \hat{\chi}(X) \geq \chi(X)$.

A.2 Iteration vector, iteration domain, and access sizes

We now prove Lemma 3:

Lemma 3. *Given the ranges of all iteration variables D^t , $t = 1, \dots, l$ during subcomputation H , if $|H| = \prod_{t=1}^l |D^t|$, then $\forall j = 1, \dots, m : |A_j(D)| = \prod_{k=1}^{\dim(\phi_j)} |D_j^k|$ and $|H|$ is maximized among all valid subcomputations which iterate over $D = [D^1, \dots, D^l]$.*

Intuition. Lemma 3 states that if each iteration variable ψ^t , $t = 1, \dots, l$ takes $|R_h^t|$ different values, then there are at most $\prod_{t=1}^l |D^t|$ different iteration vectors ψ which can be formed in H . So, intuitively, to maximize $|H|$, all combinations of values ψ^t should be evaluated. On the other hand, this also implies maximization of all access sizes $|A_j(D)| = \prod_{k=1}^{\dim(\phi_j)} |D_j^k|$.

To prove it, we now introduce two auxiliary lemmas:

Lemma 9. *For statement S , the size $|H|$ of subcomputation H (number of vertices of S computed during H) is bounded by the sizes of the iteration variables' sets R_h^t , $t = 1, \dots, l$:*

$$|H| \leq \prod_{t=1}^l |D^t|. \quad (5)$$

PROOF. Inequality 5 follows from a combinatorial argument: each computation in H is uniquely defined by its iteration vector $[\psi^1, \dots, \psi^l]$. As each iteration variable ψ^t takes $|R_h^t|$ different values during H , we have $|R_h^1| \cdot |R_h^2| \cdot \dots \cdot |R_h^l| = \prod_{t=1}^l |D^t|$ ways how to uniquely choose the iteration vector in H . \square

Now, given D , we want to assess how many different vertices are accessed for each input array A_j . Recall that this number is denoted as access size $|A_j(D)|$.

We will apply the same combinatorial reasoning to $A_j(D)$. For each access $A_j[\phi_j(\psi)]$, each one of ψ_j^k , $k = 1, \dots, \dim(\phi_j)$ iteration variables loops over set $R_{h,j}^k$ during subcomputation H . We can thus bound size of $A_j(D)$ similarly to Lemma 9:

Lemma 10. *The access size $|A_j(D)|$ of subcomputation H (the number of vertices from the array A_j required to compute H) is bounded by the sizes of $\dim(\phi_j)$ iteration variables' sets $R_{h,j}^k$, $k = 1, \dots, \dim(\phi_j)$:*

$$\forall j=1, \dots, m : |A_j(D)| \leq \prod_{k=1}^{\dim(\phi_j)} |D_j^k| \quad (6)$$

where $D_j^k \ni \psi_j^k$ is the set over which iteration variable ψ_j^k iterates during H .

PROOF. We use the same combinatorial argument as in Lemma 9. Each vertex in $A_j(D)$ is uniquely defined by $[\psi_j^1, \dots, \psi_j^{\dim(\phi_j)}]$. Knowing the number of different values each ψ_j^k takes, we bound the number of different access vectors $\phi_j(\psi_h)$. \square

Example: Consider once more statement $S1$ from LU factorization in Figure 3. We have $\phi_0 = [i, k]$, $\phi_1 = [i, k]$, and $\phi_2 = [k, k]$. Denote the iteration subdomain for subcomputation H as $D = \{[k^1, i^1], \dots, [k^{|H|}, i^{|H|}]\}$, where each variable k and i iterates over its set $k^g \in \{\psi_{k,1}, \dots, \psi_{k,K}\} = R_h^k$ and $i^g \in \{\psi_{i,1}, \dots, \psi_{i,I}\} = R_h^i$, for $g = 1, \dots, |H|$. Denote the sizes of these sets as $|R_h^k| = K_h$ and $|R_h^i| = I_h$, that is, during H , variable k takes K different values and i takes I_h different values. For ϕ_1 , both iteration variables used are different: k and i . Therefore, we have (Equation 6) $|A_1(D)| \leq K_h \cdot I_h$. On the other hand, for ϕ_2 , the iteration variable k is used twice. Recall that the access dimension is the minimum number of different iteration variables that uniquely address it (Section 2.2), so its dimension is $\dim(A_2) = 1$ and the only iteration variable needed to uniquely determine ϕ_2 is k . Therefore, $|A_2(D)| \leq K_h$.

Dominator set. Input vertices A_1, \dots, A_m form a dominator set of vertices A_0 , because any path from graph inputs to any vertex in A_0 must include at least one vertex from A_1, \dots, A_m . This is also the minimum dominator set, because of the disjoint access property (Section 2.2): any path from graph inputs to any vertex in A_0 can include at most one vertex from A_1, \dots, A_m .

Proof of Lemma 3. For subcomputation H , we have $|\bigcup_{j=1}^m A_j(D)| \leq X$ (by the definition of an X -partition). Again, by the disjoint access property, we have $\forall j_1 \neq j_2 : A_{j_1}(D) \cap A_{j_2}(D) = \emptyset$. Therefore, we also have $|\bigcup_{j=1}^m A_j(D)| = \sum_{j=1}^m |A_j(D)|$. We now want to maximize $|H|$, that is to find H_{max} to obtain computational intensity ρ (Lemma 2).

Now we prove that to maximize $|H|$, inequalities 5 and 6 must be tight (become equalities).

From proof of Lemma 9 it follows that $|H|$ is maximized when iteration vector ψ takes all possible combinations of iteration variables $\psi_h^t \in R_h^t$ during H . But, as we visit each combination of all l iteration variables, for each access A_j every combination of its $[\psi_j^1, \dots, \psi_j^{dim(\phi_j)}]$ iteration variables is also visited. Therefore, for every $j = 1, \dots, m$, each access size $|A_j(\mathcal{D})|$ is maximized (Lemma 10), as access functions are injective, which implies that for each combination of $[\psi_j^1, \dots, \psi_j^{dim(\phi_j)}]$, there is one access to A_j . $\prod_{t=1}^l |R_h^t|$ is then the upper bound on $|H|$, and its tightness implies that all bounds on access sizes $|A_j(\mathcal{D})| \leq \prod_{k=1}^{dim(\phi_j)} |D_j^k|$ are also tight. \square

A.3 Computational intensity and out-degree one vertices

Here we present a short proof of Lemma 4 followed by an example:

Lemma 4. *If in a cDAG $G = (V, E)$ every non-input vertex has at least u direct predecessors with out-degree one that are graph inputs, then the maximum computational intensity ρ of this cDAG is bounded by $\rho \leq \frac{1}{u}$.*

PROOF. By the definition of the red-blue pebble game, all inputs start in slow memory, and therefore, have to be loaded. By the assumption on the cDAG, to compute any non-input vertex $v \in V$, at least u input vertices need to have red pebbles placed on them using a load operation. Because these vertices do not have any other direct successors (their out-degree is 1), they cannot be used to compute any other non-input vertex w . Therefore, each computation of a non-input vertex requires at least u unique input vertices to be loaded. \square

Example: Consider Figure 11. In a), each compute vertex $C[i, j]$ has two input vertices: $A[i, j]$ with out-degree 1, and $b[j]$ with out-degree n , thus $u = 1$. As both array A and vector b start in the slow memory (having blue pebbles on each vertex), for each computed vertex from C , at least one vertex from A has to be loaded, therefore $\rho \leq 1$. In b), each computation needs two out-degree 1 vertices, one from vector a and one from vector b , resulting in $u = 2$. Thus, $\rho \leq \frac{1}{2}$.

A.4 Data Reuse Across Multiple Statements

Lemma 5. *The I/O cost of a program containing statements S and T which share the input array A_i is bounded by*

$$Q_{tot} \geq Q_S + Q_T - Reuse(A_i)$$

where Q_S, Q_T are the I/O costs of a program containing only statement S or T , respectively. Furthermore, we have:

$$Reuse(A_i) = \min\{|A_i(R_S)|, |A_i(R_T)|\}$$

where $|A_i(R_S)|$ and $|A_i(R_T)|$ are the total number of accesses to array A_i during the optimal execution of statements S or T separately.

PROOF. Consider an optimal sequential schedule of a cDAG G_S containing statement S only. For any subcomputation H_s and its associated iteration domain R_s its minimum dominator set is $Dom(H_s) = \bigcup_{j=1}^m A_j(R_s)$. To compute H_s , at least $\sum_{j=1}^m |A_j(R_s)| - M$ vertices have to be loaded, as only M vertices can be reused from previous subcomputations.




Figure 11: cDAGs with out-degree 1 input vertices. a) $u_a = 1, \rho_a \leq 1$. b) $u_b = 2, \rho_b \leq \frac{1}{2}$.

We seek if any loads can be avoided in the common schedule if we add statement T , denoting its cDAG G_{S+T} . Consider a subset $A_i(R_x)$ of vertices in A_i .

Consider some subset of vertices in A_i which potentially could be reused and denote it Θ_i . Now denote all vertices in A_0 (statement S) which depend on any vertex from Θ_i as Θ_S , and, analogously, set Θ_T for statement T . Now consider these two subsets Θ_S and Θ_T separately. If Θ_S is computed before Θ_T , then it had to load all vertices from Θ_i , avoiding no loads compared to the schedule of G_S only. Now, computation of Θ_T may take benefit of some vertices from Θ_i , which can still reside in fast memory, avoiding up to $|\Theta_i|$ loads.

The total number of avoided loads is bounded by the number of loads from A_i which are shared by both S and T . Because statement S loads at most $|A_i(R_S)|$ vertices from A_i during optimal schedule of G_S , and T loads at most $|A_i(R_T)|$ of them for G_T , the upper bound of shared, and possibly avoided loads is $Reuse(A_i) = \min\{|A_i(R_S)|, |A_i(R_T)|\}$. \square

The **reuse size** is defined as $Reuse(A_i) = \min\{|A_i(R_S)|, |A_i(R_T)|\}$. Now, how to find $|A_i(R_S)|$ and $|A_i(R_T)|$?

Observe that $|A_i(R_S)|$ is a property of G_S , that is, the cDAG containing statement S only. Denote the I/O optimal schedule parameters of G_S : V_{max}^S, X_0^S , and $|A_i(R_{max}^S(X_0^S))|$ (Section 3.2). Similarly, for G_T : V_{max}^T, X_0^T , and $|A_i(R_{max}^T(X_0^T))|$. We now derive: 1) at least how many subcomputations does the optimal schedule have: $s \geq \frac{|V|}{|H_{max}|}$, 2) at least how many accesses to A_i are performed per optimal subcomputation $|A_i(R_{max}(X_0))|$. Then:

$$Reuse(A_i) = \min\{|A_i(R_{max}^S(X_0^S))| \frac{|V^S|}{|V_{max}^S|}, |A_i(R_{max}^T(X_0^T))| \frac{|V^T|}{|V_{max}^T|}\} \quad (7)$$

We now proceed to Lemma 6

Lemma 6. *Any dominator set of set $B_j(\mathcal{D})$ must be of size at least*

$$|Dom(B_j(\mathcal{D}))| \geq \frac{|B_j(\mathcal{D})|}{\rho_s}$$

PROOF. By Lemma 1, for one loaded vertex, we may compute at most ρ_s vertices of A_0 . These are also vertices of B_j . Thus, to

compute $|B_j(\mathcal{D})|$ vertices of B_j , at least $\frac{|B_j(\mathcal{D})|}{\rho_S}$ loads must be performed. We just need to show that at least that many vertices have to be in any dominator set $\text{Dom}(B_j(\mathcal{D}))$. Now, consider the converse: There is a vertex set $D = \text{Dom}(B_j(\mathcal{D}))$ such that $|D| < \frac{|B_j(\mathcal{D})|}{\rho_S}$. But that would mean, that we could potentially compute all $|B_j(\mathcal{D})|$ vertices by only loading $|D|$ vertices, violating Lemma 1. \square

A.5 Parallel I/O Lower Bounds

Lemma 7. *The minimum number of I/O operations in a parallel pebble game, played on a cDAG with $|V|$ vertices with P processors each equipped with M pebbles, is $Q \geq \frac{|V|}{P \cdot \rho}$, where ρ is the maximum computational intensity independent of P (Lemma 1).*

PROOF. Following the analysis of Section 3 and the parallel machine model (Section 5), the computational intensity ρ is independent of a number of parallel processors - it is solely a property of a cDAG and private fast memory size M . Therefore, following Lemma 1, what changes with P is the volume of computation $|V|$, as now at least one processor will compute at least $|V_p| = \frac{|V|}{P}$ vertices. By the definition of the computational intensity, the minimum number of I/O operations required to pebble these $|V_p|$ vertices is $\frac{|V_p|}{\rho}$. \square